Week 4 midpoint, design proposal

Jarvaids



The wall shown is a drawing made for the demo, not a photograph of a session

A meeting assistant that records what is said and what is put on the wall during a design session, while the session is running.

Group 3 (tentative)
Luuk Brok, Pieter Borst, Julie van Veelen, Roselique van den Brink

Project background

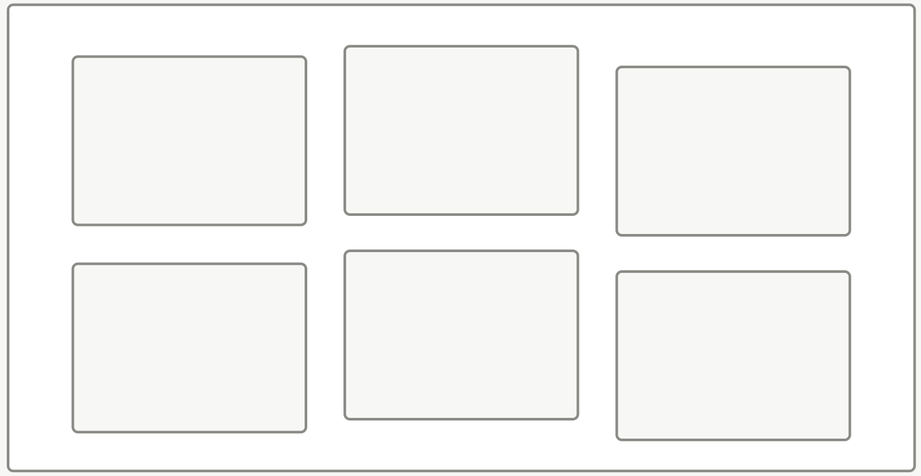
Design teams in working sessions produce remarks, sketches and sticky notes. Most of it is not recorded. Three problems recur.

1. A remark made earlier in the session cannot be found again.
2. The wall is photographed at the end, but the notes are rarely typed up.
3. The board survives the meeting; the reasoning behind it does not.

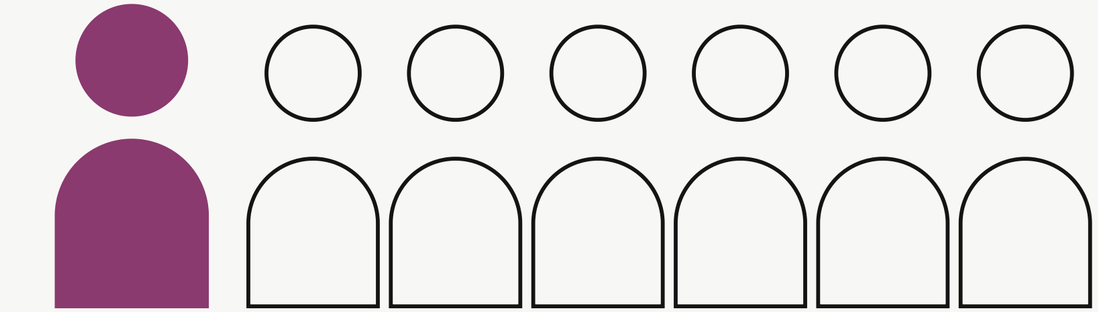
Source: concept 001, ideation folder

Users and context

Designers, technical participants in multidisciplinary teams, and Master’s students in design programmes, in meetings and brainstorm sessions that combine spoken discussion, sketches, sticky notes and photographs. Two to eight people, one to two hours, one facilitator, a shared screen.



The wall



Roles: flowchart F4.3; session size: app/server/README.md

Product concept

An AI tool used during design meetings and brainstorm sessions. It takes the spoken discussion and photographs of the room as input, and keeps a structured visual overview of the meeting on the digital whiteboard while the meeting is running. Participants can ask it questions and give it commands by voice.

Input	Speech in the room	Photographs of the wall, sketches and screens	Typed notes	Voice commands
Overview	Themes and clusters	Decisions and actions	Open questions	Notes from the wall
On request	Answers about the meeting so far, with the moment cited		Suggestions for next steps, marked as the tool’s	

Source: the structured brainstorm notes of 11 September and concept 001

Features

Live transcription of the discussion, with speaker labels.

F1.2, F1.3. Plumbing

Voice commands, such as “summarise this” or “add this to the planning”, carried out on the whiteboard.

F2.5, F2.6. Proposed

Photographs of the sticky-note wall become individual notes on the whiteboard.

F1.4 to F1.6. Built with a placeholder reader

A live overview of themes, decisions, actions and open questions, laid out rather than listed.

F2.1 to F2.4, F3.2, F3.3. Proposed

Questions about the meeting so far, answered live and citing the moment the answer came from.

F2.7, F6.3. Proposed

Anything the tool adds is provisional until a participant accepts it. Corrections are recorded.

F3.4, F3.5, F6.5. Built

Consent before capture, audio deleted after the session, transcripts anonymised.

F6.1, F6.2, F6.6. Built, audio proposed

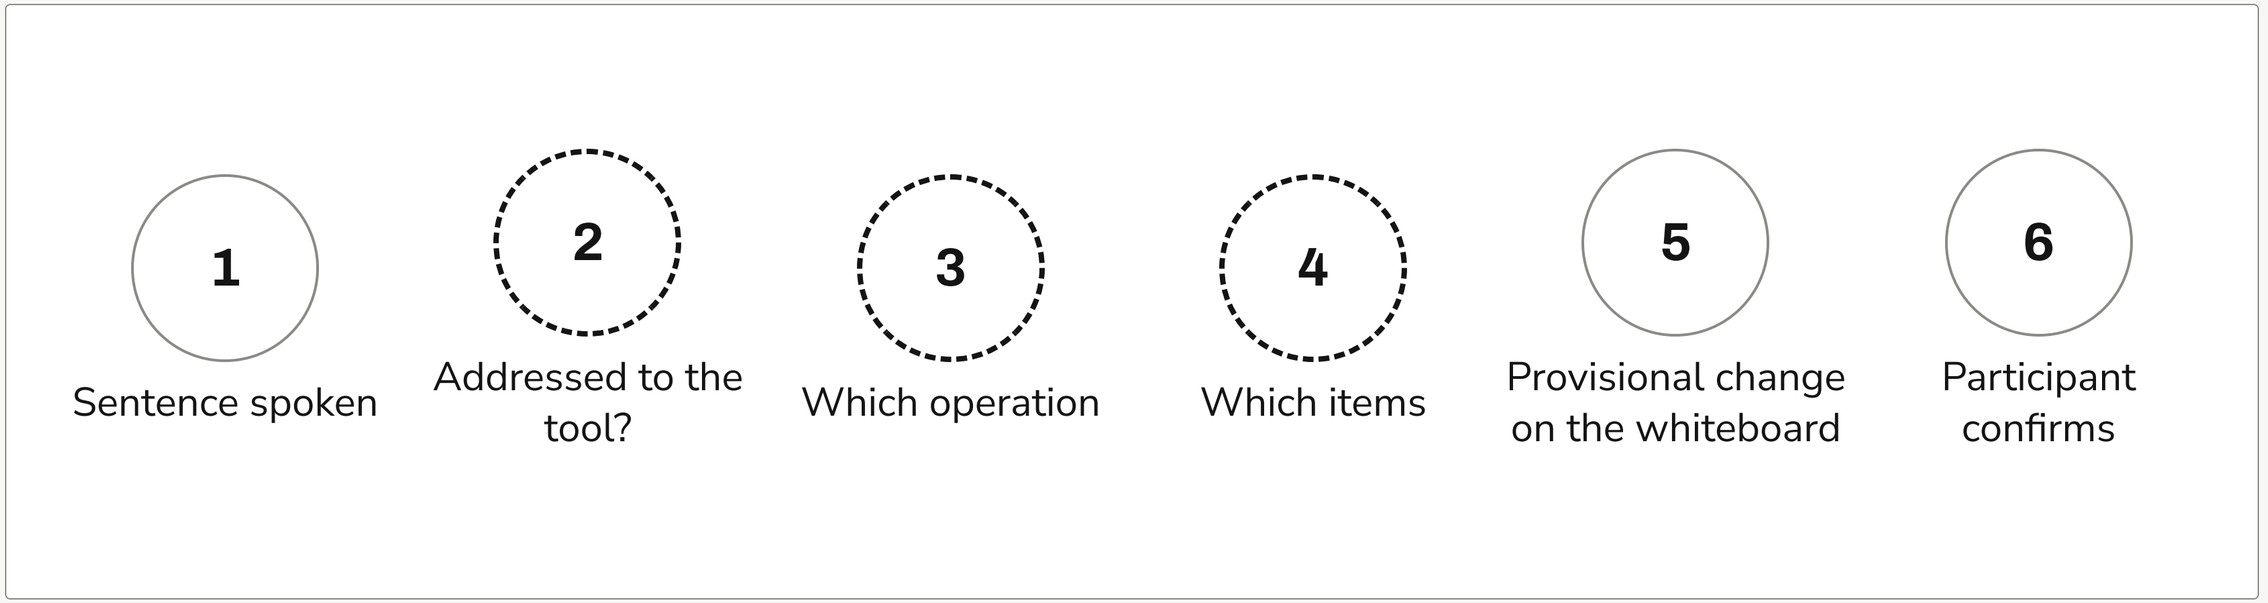
Export of the canvas and the minutes; push to Miro or FigJam.

F5.1 to F5.3. Export built, push proposed

F-numbers: documentation/flowchart.md. Status of each function: documentation/proposal.md

Voice commands

A participant says “add this to the planning”. The command interpreter decides whether the sentence was addressed to the tool, works out what “this” refers to from the canvas and the last minute of speech, and proposes one change to the whiteboard. A participant confirms it. When the reference is ambiguous the tool asks rather than guesses.



Pipeline: documentation/pipelines.md section 3. Dashed stages involve a model. Status: proposed

Findings from week 3

Generic AI feedback was found unhelpful. Feedback on a specific element was useful.

Group reflection

Five simulated interviewees responded positively; five real interviewees responded negatively.

Interviews, via the MapIt report

The MapIt report presented one interviewee's view as a conclusion of the group.

MapIt report

The group asked for a live overview in which AI suggestions are visibly distinct from what was said.

FrameIt conversation

Annotated material: deliverables/week4-midpoint/annotations

Design question

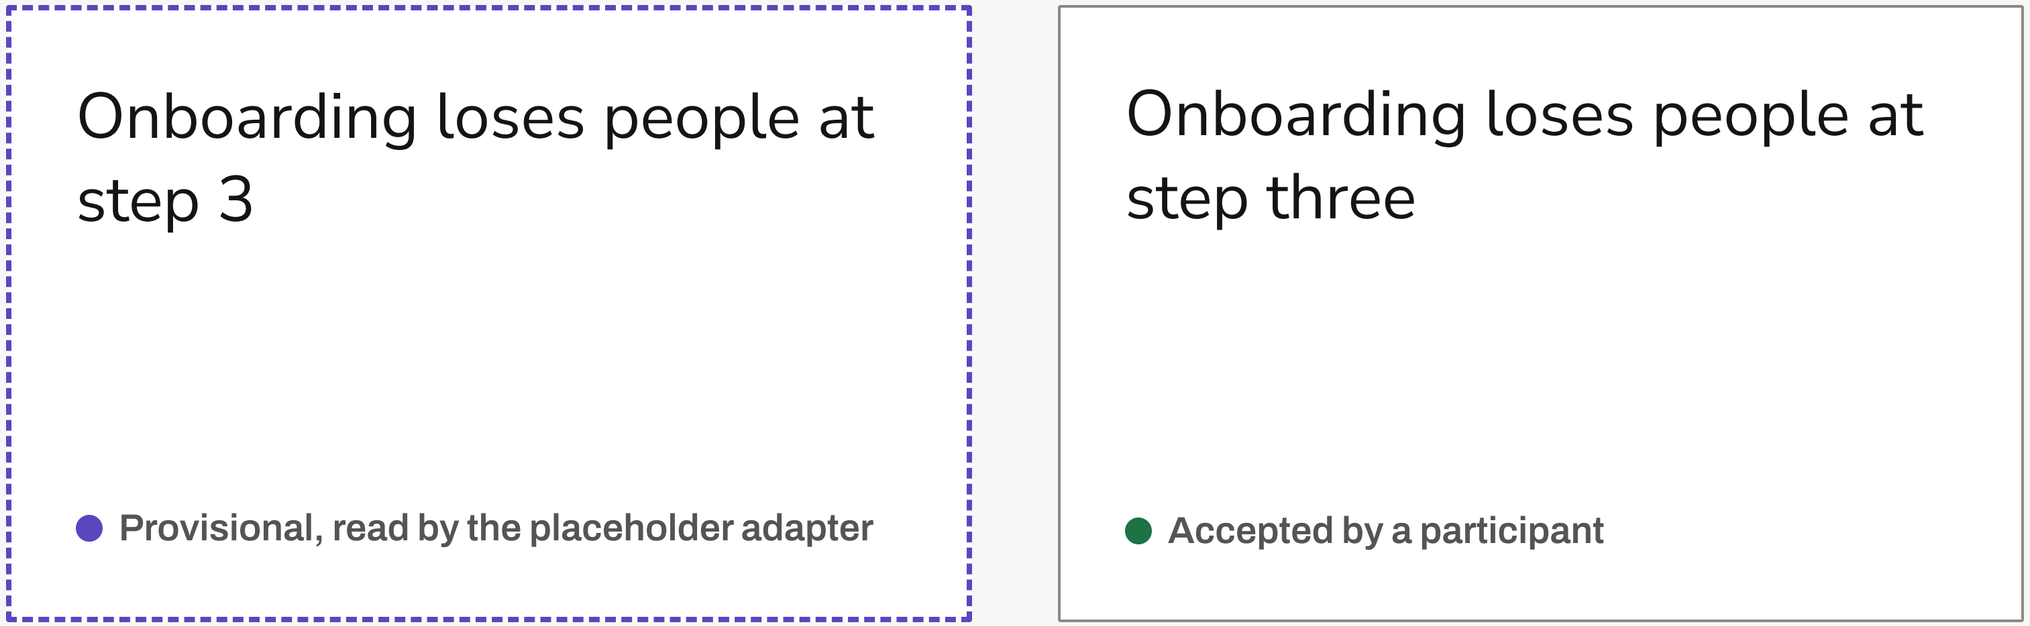
> How can what is said and written during a design session be recorded as it happens, in a way that keeps the team’s own work distinct from what a model proposed?

Derived from the concept, the week 3 findings and the flowchart

Design principle

Provisional output

Anything produced by a model is marked provisional until a participant accepts it. A provisional note has a dashed edge and names its source. Moving a note does not count as accepting it.



Flowchart F3.4 and F3.5; design system rules 16 to 18

User journey

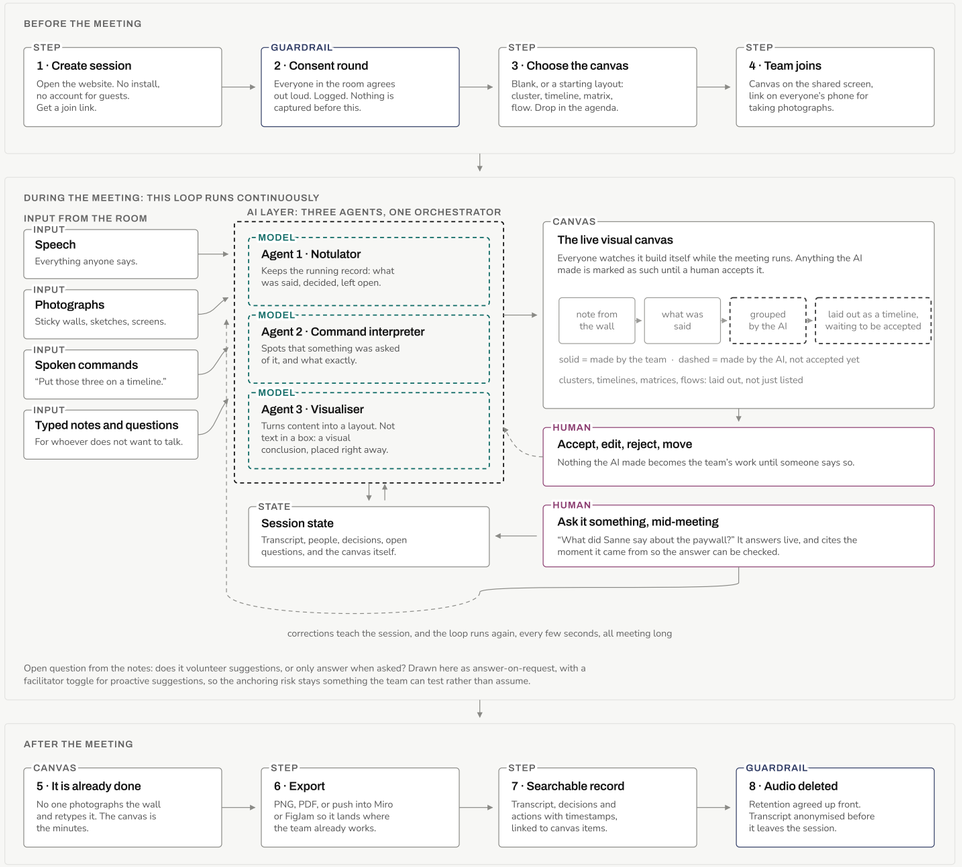
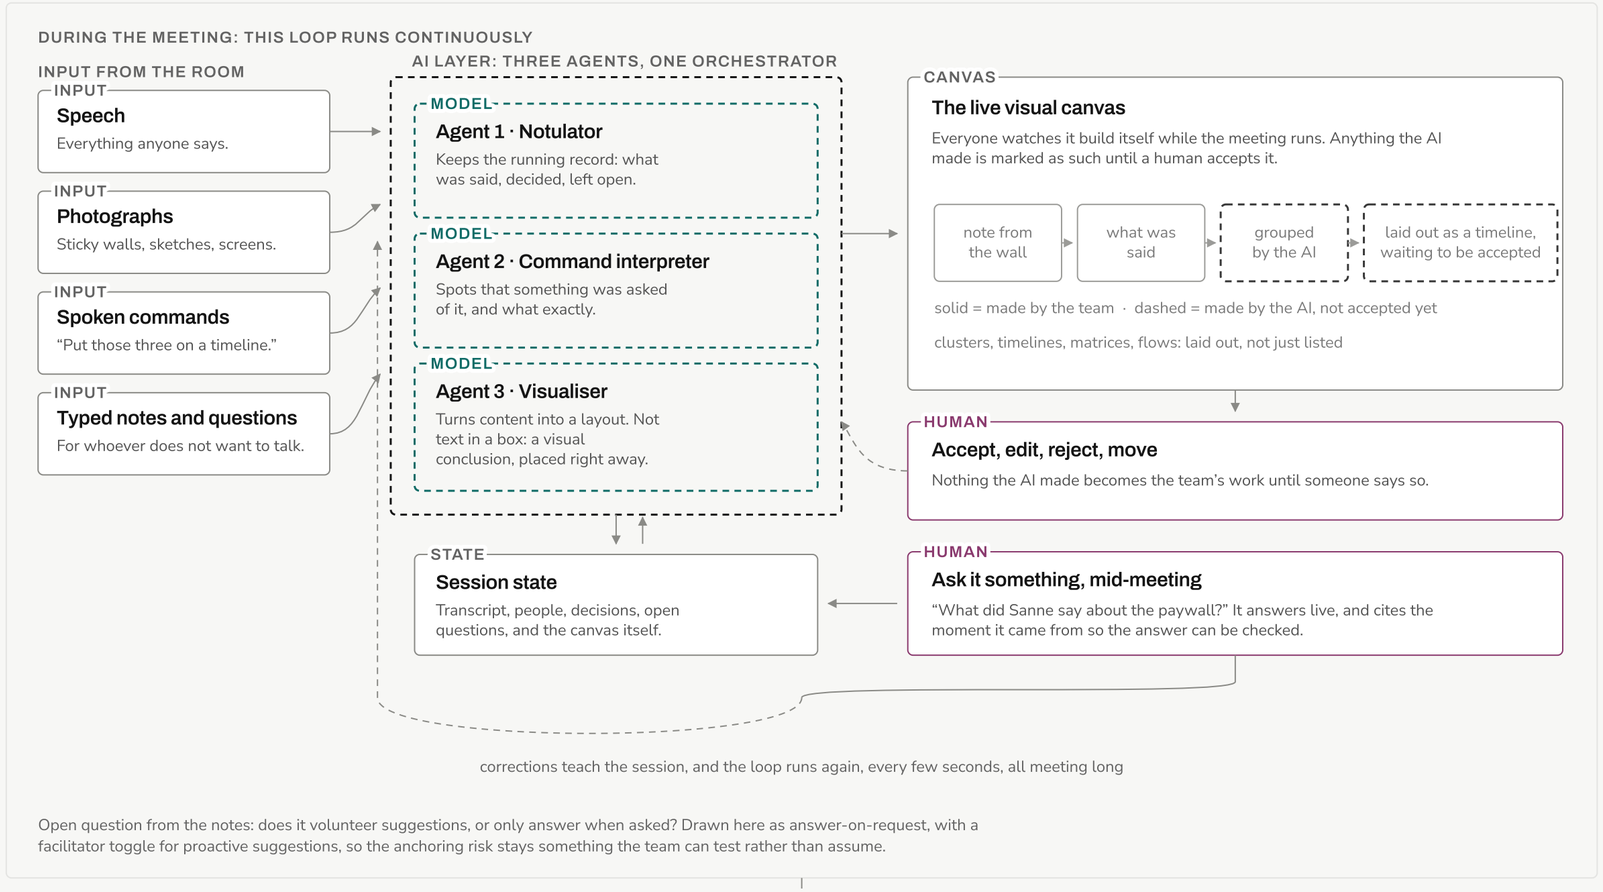


Diagram: documentation/flowchart.html, user journey tab. Dashed boxes are where a model runs.

User journey, during the meeting



The middle band of the previous diagram. Input on the left, agents in the middle, the canvas and the participants on the right.

System architecture

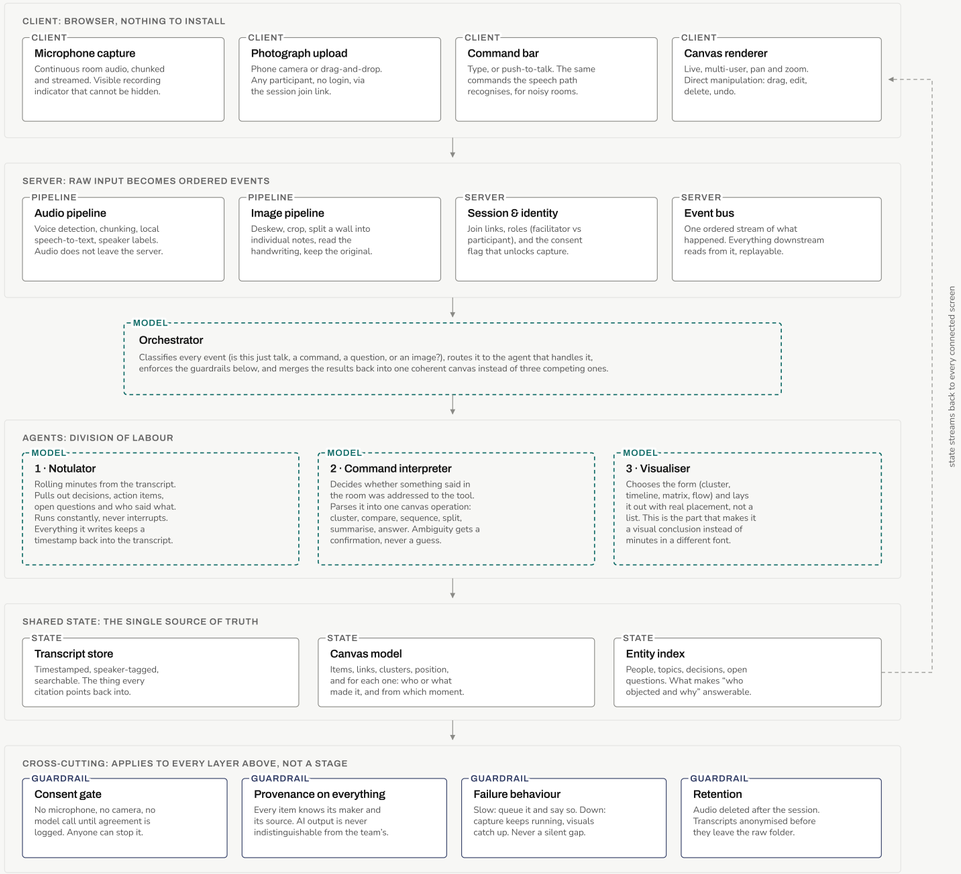


Diagram: documentation/flowchart.html, system architecture tab. Agent boxes describe responsibilities, not separate services.

Functions in scope

The flowchart defines thirty-seven functions in six groups. Each one has a status: built, plumbing (the contract exists, the model does not), proposed, or open.

7
F1 Capture


8
F2 Understanding


7
F3 Visualisation


5
F4 Session

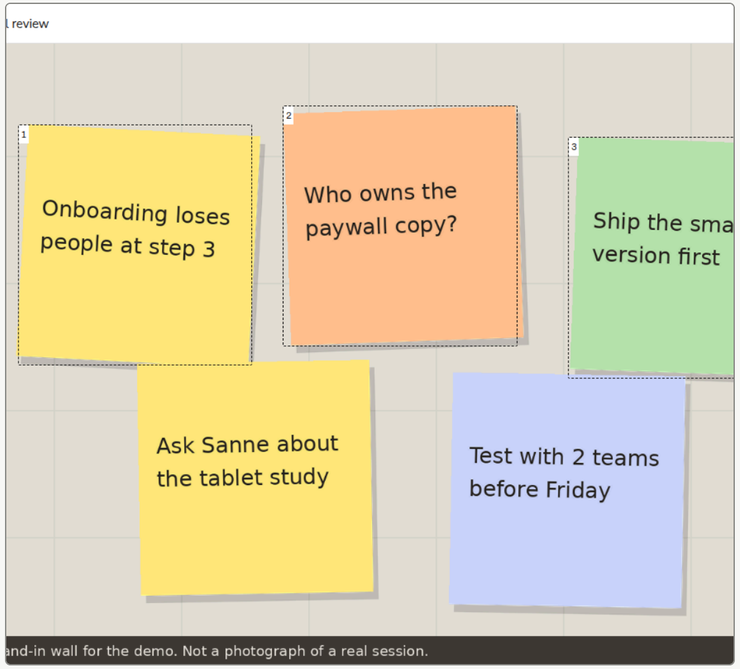

4
F5 Output


6
F6 Trust

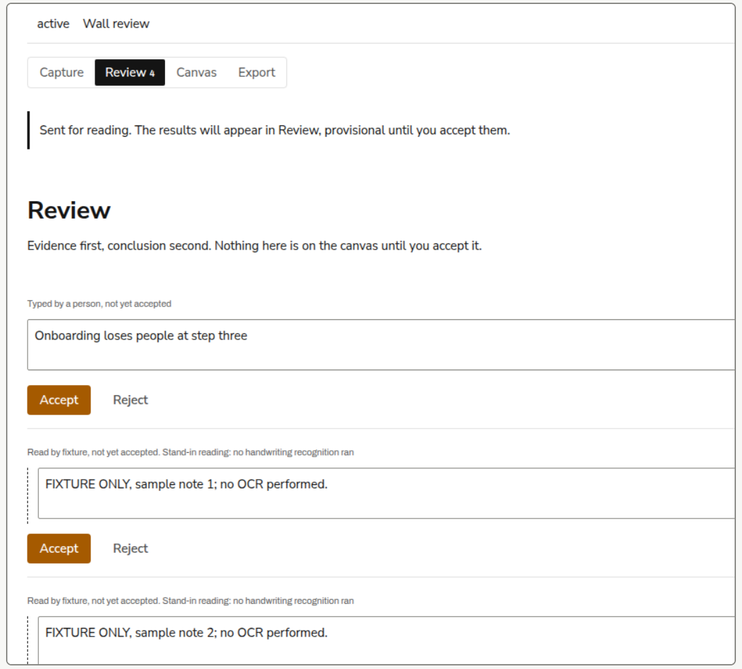

Filled: built. Half: plumbing. Dashed: proposed. Empty: open. Full table in documentation/proposal.md

From photograph to notes

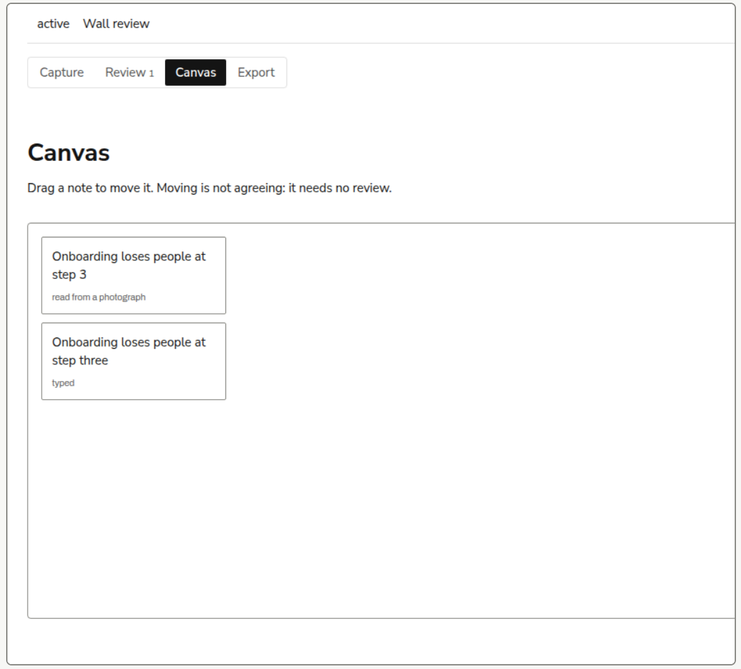
A participant marks each note on the photograph. A placeholder adapter returns a reading; no handwriting recognition has been run. A participant corrects, accepts or rejects each reading.



Notes are marked



Readings arrive for review

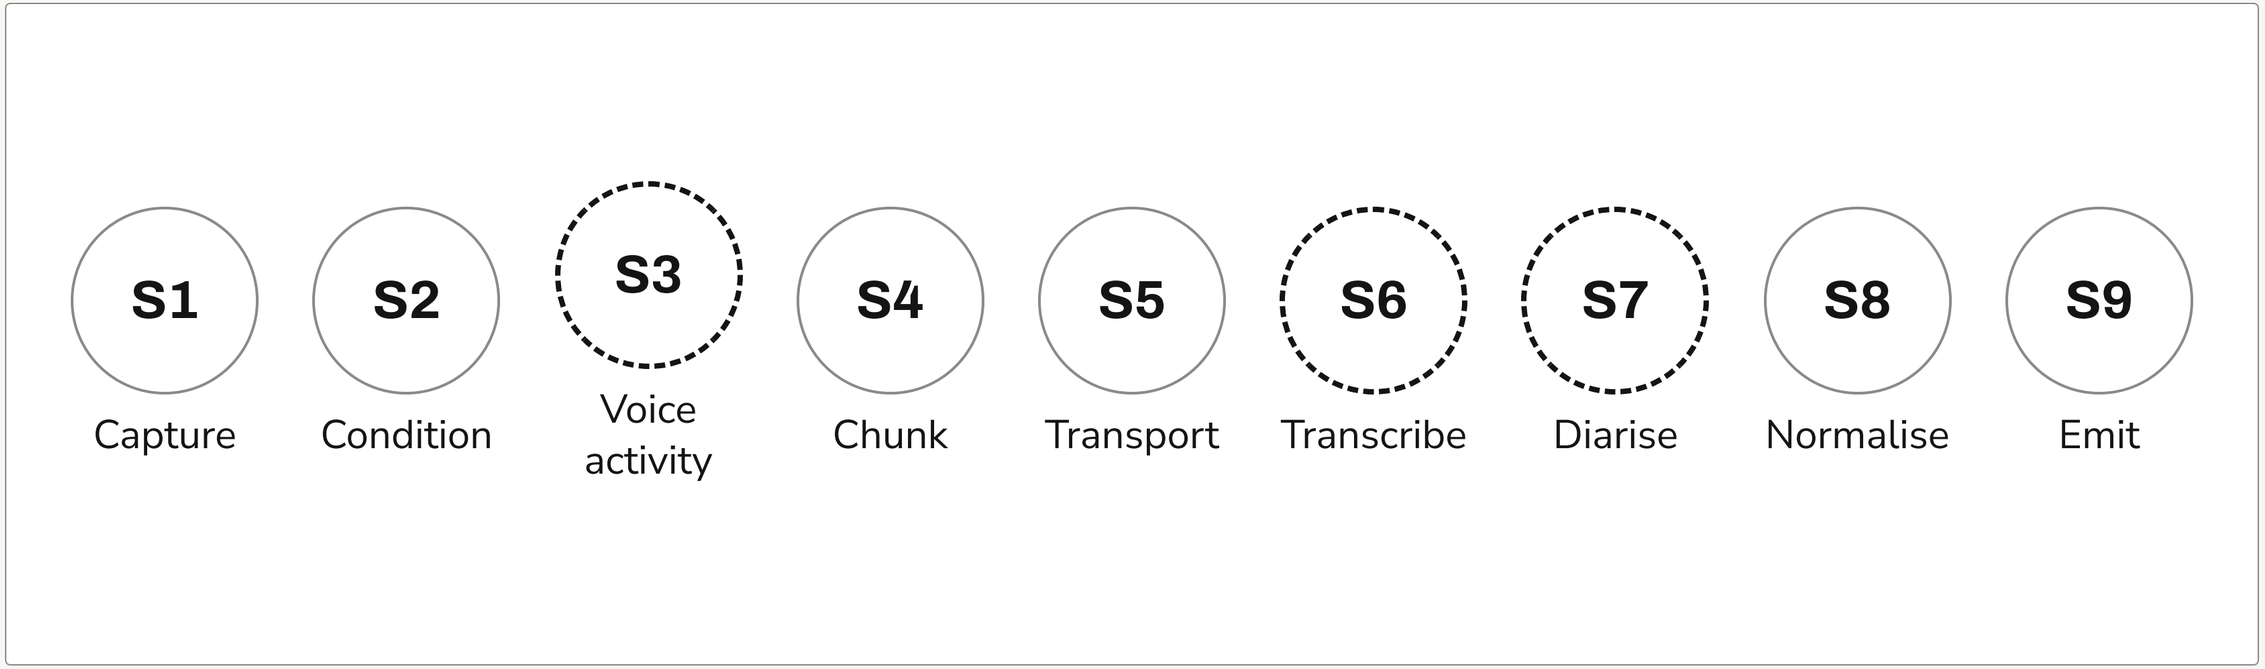


Accepted notes reach the canvas

Captures of the client. Proposed automatic segmentation and reading: documentation/stack.md

Speech pipeline

Nine stages are designed; three involve a model. A server contract, a browser chunker and a placeholder transcriber exist. The wrapper for a local Whisper model is written and has not been run. No speech recognition has been run.



Stages: documentation/pipelines.md section 1. Status: app/server/BUILD-STATUS.md

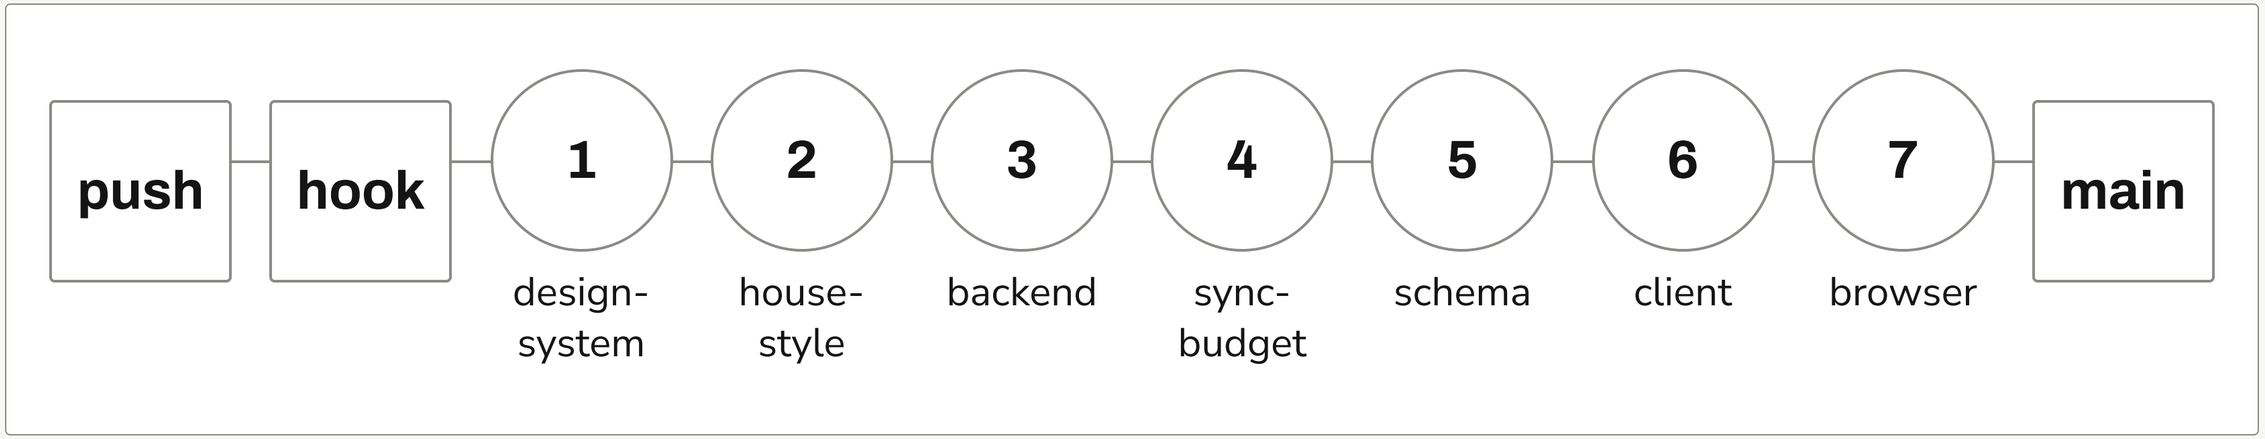
Technology stack

Client	ES modules, no build step	capture-audio.mjs	d3-force
Transport	HTTP and JSON	websockets	
Server	Python, SQLite	Pillow	OpenCV
Adapters	FixtureExtractor	FixtureTranscriber	LocalWhisperTranscriber
Consent gate			
Hosted	Claude, GPT-4o, Gemini	Miro REST API v2	

Solid: in the repository. Dashed: candidate, not run. Names, licences and reasons: documentation/stack.md

Repository and verification

Every push runs seven checks before it can be merged.



254
backend tests

150
mutation guards

13
design audit checks

105
contrast pairs checked

Jobs: .github/workflows/verify.yml. Counts: AGENTS.md. Description: documentation/how-we-work.md

Live demonstration

A photographed wall is sent for reading. One reading is corrected and accepted, one is rejected, and an accepted note is moved on the canvas. The readings are placeholders; the workflow around them is real.

FIXTURE ONLY, sample note 1; no OCR performed.

● Provisional, read by the placeholder adapter

Sequence: deliverables/week4-midpoint/demo/RUNBOOK.md

Riskiest assumption and next steps

The first test is whether notes on a real wall can be read well enough to be worth correcting. The thresholds are agreed before the test. Speech capture follows, subject to consent from everyone present.

10

photographs of real walls

80%

of notes correctly separated

80%

of notes correctly read

Test and thresholds: documentation/pipelines.md, unknowns table